

Assignment 10

Prepared by:

- Josh Levine (jl2108)
- Harsh Patel (hkp49)
- Jaini Patel (jp1891)
- Yifan Liao (yl1463)
- Aayush Shah (avs93)

Sequential Probability Ratio Test

A sequential probability ratio test (SPRT) is a hypothesis test for sequential samples.

Sequential sampling works in a very non-traditional way; instead of a fixed sample size, you choose one item (or a few) at a time, and then test your hypothesis. You can either:

- Reject the null hypothesis (H_0) in favor of the alternate hypothesis (H_1) and stop,
- Keep the null hypothesis and stop,
- Fail to reach either conclusion and continue sampling.

If you fail to reach a conclusion, you repeat the sampling and then the hypothesis test. You keep on repeating this process until you have a sound conclusion, so you don't know how big your sample will be until you're finished testing.

As in classical hypothesis testing, SPRT starts with a pair of hypotheses, say H_0 and H_1 for the null hypothesis and alternative hypothesis respectively. They must be specified as follows:

$$H_0 : p = p_0 \quad H_1 : p = p_1$$

The next step is calculate the cumulative sum of the log-likelihood ratio, $\log \Lambda_i$, as new data arrive:

$$S_i = S_{i-1} + \log \Lambda_i$$

The stopping rule is a simple thresholding scheme:

$a < S_i < b$: continue monitoring (critical inequality)

$S_i \geq b$: Accept H_1

$S_i \leq a$: Accept H_0

where a and b ($0 < a < b < \infty$) are thresholds that depend on the desired type I and type II errors.

For this assignment,

S = cumulative sum of logs

$h_0 = \text{hypothesis0} = p \leq p_1$

$h_1 = \text{hypothesis1} = p \geq p_2$

$\alpha_1 = 0.01$; $\alpha_2 = 0.01$

We have our SPRT function that calculates our log sum S and then checks which hypothesis, h_0 or h_1 , is correct depending on the following criteria:

$A < S < B$ - continue loop

$S \geq B$ - h_1 accepted

$S \leq A$ - h_0 accepted

where A and B are the thresholds we get based on α_1 and α_2 errors.

SPRT then returns the correct value to simulate. Then the simulate function (which calls SPRT) runs 100 iterations to choose either h_0 or h_1 depending on the higher count of accepted hypothesis and also calculates average steps to convergence.

Below is the code for SPRT function:

```
# SPRT function

sprt <- function(alp1 = 0.01, alp2 = 0.01,
                 p1 = 0.45, p2 = 0.55, bernoulli_p = 0.3)
{
  S = 0
  log_test = 0
  converge_count = 0

  # calculating stop threshold
  A = log(alp2/(1 - alp1))
  B = log((1 - alp2)/alp1)
  valid_hypothesis = -1

  while(TRUE){

    converge_count = converge_count + 1

    # generate bernoulli RV
    data_point = rbinom(1, 1, bernoulli_p)

    # log ratio
    log_test = (data_point*p2 + (1-data_point)*(1-p2))
    - (data_point*p1 + (1-data_point)*(1-p1))

    # cumulative sum
    S = S + log_test

    # check stop conditions
    if(S>=B){
      #Accept H1
      valid_hypothesis = 1
      break
    }
    if(S<=A){
      #Accept H0
      valid_hypothesis = 0
      break
    }
    #print(paste("data_point", data_point, "log_test", log_test, "S", S))
  }
}
```

```

    return(list(converge_count = converge_count, valid_hypothesis = valid_hypothesis))
}

```

The following code shows the simulation function:

```

#Simulation Function
simulate_sprt <- function(bernoulli_p = 0.3, iterations = 100){

  sum = 0
  h0 = 0
  h1 = 0

  # Averaging SPRT
  for(i in c(1:iterations)){
    # call the SPRT function that generates bernoulli RV with bernoulli_p
    result = sprt(bernoulli_p=bernoulli_p)

    # H0 is accepted
    if(result$valid_hypothesis == 0){
      h0 = h0 + 1
    }
    # H1 is true
    if(result$valid_hypothesis == 1){
      h1 = h1 + 1
    }
    sum = sum + result$converge_count
  }
  average = sum/iterations
  # print(paste("bernoulli_p=", bernoulli_p, ", avg_n=", average))
  return(list(average=average, h0=h0, h1=h1))
}

```

Below is the code to run the simulation function for a range of p values:

```

# Running the simulation function for a range of p values
test_list = seq(0.3, 0.68, by=0.02)

# data frame to store the result
img<-data.frame(bernoulli_p=numeric(),
                avg_steps_to_converge=numeric(), h0=numeric(),
                h1=numeric(), hypothesis_accepted=character())

#iterating over test_list
for(p in test_list){
  print(p)
  result = simulate_sprt(bernoulli_p = p)

  if(result$h0 < result$h1)
    accepted_h = "H1"
  if(result$h0 > result$h1 ){
    accepted_h = "H0"
  }
  if(result$h1 == result$h0){
    accepted_h = "H0/H1"
  }
}

```

```

new_data = data.frame(bernolli_p=p,
                      avg_steps_to_converge=result$average,
                      h0=result$h0, h1=result$h1,
                      accepted_hypothesis=accepted_h,
                      stringsAsFactors = FALSE)
img<-rbind(img, new_data)
}

```

```

## [1] 0.3
## [1] 0.32
## [1] 0.34
## [1] 0.36
## [1] 0.38
## [1] 0.4
## [1] 0.42
## [1] 0.44
## [1] 0.46
## [1] 0.48
## [1] 0.5
## [1] 0.52
## [1] 0.54
## [1] 0.56
## [1] 0.58
## [1] 0.6
## [1] 0.62
## [1] 0.64
## [1] 0.66
## [1] 0.68

```

To show our results the code displays a table run on a sequence of Bernoulli from 0.3 to 0.68 with gap of 0.02

```

#display img
img

```

	bernolli_p	avg_steps_to_converge	h0	h1	accepted_hypothesis
## 1	0.30	9.95	0	100	H1
## 2	0.32	9.98	0	100	H1
## 3	0.34	9.97	0	100	H1
## 4	0.36	9.96	0	100	H1
## 5	0.38	9.94	0	100	H1
## 6	0.40	9.95	0	100	H1
## 7	0.42	9.86	0	100	H1
## 8	0.44	9.82	0	100	H1
## 9	0.46	9.81	0	100	H1
## 10	0.48	9.79	0	100	H1
## 11	0.50	9.77	0	100	H1
## 12	0.52	9.69	0	100	H1
## 13	0.54	9.64	0	100	H1
## 14	0.56	9.59	0	100	H1
## 15	0.58	9.51	0	100	H1
## 16	0.60	9.50	0	100	H1
## 17	0.62	9.39	0	100	H1
## 18	0.64	9.41	0	100	H1
## 19	0.66	9.33	0	100	H1
## 20	0.68	9.30	0	100	H1

Conclusion

* Why does it give the result you got?

The results in general show that p values from 0.45 to 0.55 require the highest number of iterations to converge with the maximum at $p = 0.5$. This is because p is equidistant from both $p = 0.45$ and $p = 0.55$ so the log sum S will keep changing and going from one hypothesis to other until eventually it converges to either h_1 or h_0 based on the count. Both h_0 and h_1 are accepted almost equal amount of times at $p = 0.5$ so we cannot know for sure which will we see until the end.

As we go farther out from $p = 0.45$ and $p = 0.55$ we see that it takes less and less iterations to converge as S just goes to either h_0 or h_1 fairly soon. So, For $p < 0.5$ h_0 is accepted $p > 0.5$ h_1 is accepted and at $p = 0.5$ both of equal likelihood of getting accepted.

* What do you think it would do for .54 ? Try it.

Now for $p = 0.54$. From our above conclusions we should see that at $p = 0.54$ h_1 should be accepted because there is a higher probability for h_1 to accepted (h_0 is rejected) as it is closer to 0.55 than h_0 . Also, $p = 0.54$ should have lower steps than $p = 0.5$ to converge. Both of these results can be verified based on our output table. Hence SPRT function works as expected.